

Load Dataset

```
import pandas as pd

data = pd.read_csv('WineQT.csv')
```

Explore Dataset

```
print("Dataset Info:")
print(data.info())
print("\nFirst 5 Rows:")
print(data.head())
print("\nDescriptive Statistics:")
print(data.describe())
```

```
11  quality          1143 non-null   int64
12  Id              1143 non-null   int64
dtypes: float64(11), int64(2)
memory usage: 116.2 KB
None
```

First 5 Rows:

	s	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.4	0.70	0.00	1.9	0.076	
1	7.8	0.88	0.00	2.6	0.098	
2	7.8	0.76	0.04	2.3	0.092	
3	11.2	0.28	0.56	1.9	0.075	
4	7.4	0.70	0.00	1.9	0.076	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
0	11.0	34.0	0.9978	3.51	0.56	
1	25.0	67.0	0.9968	3.20	0.68	
2	15.0	54.0	0.9970	3.26	0.65	
3	17.0	60.0	0.9980	3.16	0.58	
4	11.0	34.0	0.9978	3.51	0.56	

	alcohol	quality	Id
0	9.4	5	0


count	1143.000000	1143.000000	1143.000000	1143.000000
mean	0.086933	15.615486	45.914698	0.996730
std	0.047267	10.250486	32.782130	0.001925
min	0.012000	1.000000	6.000000	0.990070
25%	0.070000	7.000000	21.000000	0.995570
50%	0.079000	13.000000	37.000000	0.996680
75%	0.090000	21.000000	61.000000	0.997845
max	0.611000	68.000000	289.000000	1.003690

	pH	sulphates	alcohol	quality	Id
count	1143.000000	1143.000000	1143.000000	1143.000000	1143.000000
mean	3.311015	0.657708	10.442111	5.657043	804.969379
std	0.156664	0.170399	1.082196	0.805824	463.997116
min	2.740000	0.330000	8.400000	3.000000	0.000000
25%	3.205000	0.550000	9.500000	5.000000	411.000000
50%	3.310000	0.620000	10.200000	6.000000	794.000000
75%	3.400000	0.730000	11.100000	6.000000	1209.500000
max	4.010000	2.000000	14.900000	8.000000	1597.000000

Distribution of Target Variable ('quality')


```
import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(12, 8))
sns.heatmap(data.corr(), annot=True, cmap='coolwarm', fmt='.2f')
plt.title("Correlation Matrix")
plt.show()
```



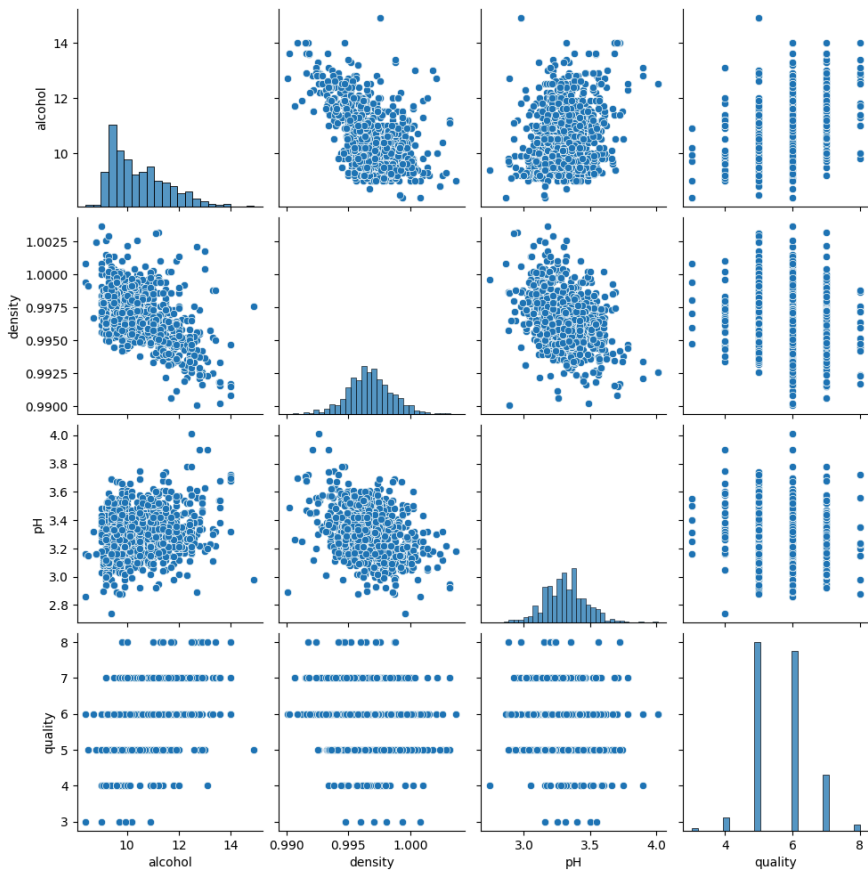
Distribution of Target Variable '(quality)'

```
plt.figure(figsize=(8, 6))
sns.countplot(x='quality', data=data)
plt.title("Distribution of Wine Quality")
plt.show()
```



Pair Plot

```
sns.pairplot(data[['alcohol', 'density', 'pH', 'quality']])  
plt.show()
```



Preprocessing Separate Features and Target

```
X = data.drop(columns=['quality'])
y = data['quality']
```

Train-Test Split

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=
```

Feature Scaling

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

Model Building and Evaluation **bold text** Random Forest Classifier

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report

rf = RandomForestClassifier(random_state=42)
rf.fit(X_train_scaled, y_train)
y_pred_rf = rf.predict(X_test_scaled)
print("\nRandom Forest Accuracy:", accuracy_score(y_test, y_pred_rf))
print("Classification Report:\n", classification_report(y_test, y_pred_rf))
```



Random Forest Accuracy: 0.6899563318777293

Classification Report:

	precision	recall	f1-score	support
4	0.00	0.00	0.00	6
5	0.73	0.75	0.74	96
6	0.64	0.71	0.67	99
7	0.76	0.62	0.68	26
8	0.00	0.00	0.00	2
accuracy			0.69	229
macro avg	0.43	0.41	0.42	229
weighted avg	0.67	0.69	0.68	229

```
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: U
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: U
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: U
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

Stochastic Gradient Descent **Classifier**

```

from sklearn.linear_model import SGDClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report

rf = RandomForestClassifier(random_state=42)
rf.fit(X_train_scaled, y_train)
y_pred_rf = rf.predict(X_test_scaled)
print("\nRandom Forest Accuracy:", accuracy_score(y_test, y_pred_rf))
print("Classification Report:\n", classification_report(y_test, y_pred_rf))

sgd = SGDClassifier(random_state=42)
sgd.fit(X_train_scaled, y_train)
y_pred_sgd = sgd.predict(X_test_scaled)
print("\nSGD Accuracy:", accuracy_score(y_test, y_pred_sgd))
print("Classification Report:\n", classification_report(y_test, y_pred_sgd))

```



Random Forest Accuracy: 0.6899563318777293

Classification Report:

	precision	recall	f1-score	support
4	0.00	0.00	0.00	6
5	0.73	0.75	0.74	96
6	0.64	0.71	0.67	99
7	0.76	0.62	0.68	26
8	0.00	0.00	0.00	2
accuracy			0.69	229
macro avg	0.43	0.41	0.42	229
weighted avg	0.67	0.69	0.68	229

SGD Accuracy: 0.5851528384279476

Classification Report:

	precision	recall	f1-score	support
4	0.00	0.00	0.00	6
5	0.68	0.78	0.73	96
6	0.66	0.39	0.49	99
7	0.34	0.77	0.47	26
8	0.00	0.00	0.00	2
accuracy			0.59	229
macro avg	0.34	0.39	0.34	229
weighted avg	0.61	0.59	0.57	229

```

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: U
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: U
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: U
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: U
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: U
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: U
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

Support Vector **Classifier**

```

from sklearn.svm import SVC # Import the SVC class
from sklearn.metrics import accuracy_score, classification_report

svc = SVC(random_state=42)
svc.fit(X_train_scaled, y_train)
y_pred_svc = svc.predict(X_test_scaled)
print("\nSVC Accuracy:", accuracy_score(y_test, y_pred_svc))
print("Classification Report:\n", classification_report(y_test, y_pred_svc))

```



```

SVC Accuracy: 0.6550218340611353
Classification Report:

```

	precision	recall	f1-score	support
4	0.00	0.00	0.00	6
5	0.70	0.75	0.72	96
6	0.61	0.70	0.65	99
7	0.69	0.35	0.46	26
8	0.00	0.00	0.00	2
accuracy			0.66	229
macro avg	0.40	0.36	0.37	229
weighted avg	0.64	0.66	0.64	229

```

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: U
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: U
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: U
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

Hyperparameter Tuning **Random Forest** Example

```

param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5]
}

```

Best Parameters for Random Forest and **Best** Cross-Validation Score

```

from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier

```

```

param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5]
}

```

```

grid_search = GridSearchCV(RandomForestClassifier(random_state=42), param_grid, cv=5,

```



```
grid_search.fit(X_train_scaled, y_train)
print("\nBest Parameters for Random Forest:", grid_search.best_params_)
print("Best Cross-Validation Score: ", grid_search.best_score_)
```



```
Best Parameters for Random Forest: {'max_depth': 20, 'min_samples_split': 2, 'n_es':
Best Cross-Validation Score: 0.664042514862187
```

Evaluate Best Model

```
best_rf = grid_search.best_estimator_
y_pred_best = best_rf.predict(X_test_scaled)
print("\nBest Random Forest Accuracy:", accuracy_score(y_test, y_pred_best))
print("Classification Report:\n", classification_report(y_test, y_pred_best))
```



```
Best Random Forest Accuracy: 0.6899563318777293
```

```
Classification Report:
```

	precision	recall	f1-score	support
4	0.00	0.00	0.00	6
5	0.73	0.75	0.74	96
6	0.64	0.71	0.67	99
7	0.76	0.62	0.68	26
8	0.00	0.00	0.00	2
accuracy			0.69	229
macro avg	0.43	0.41	0.42	229
weighted avg	0.67	0.69	0.68	229

```
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: U
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: U
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: U
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

Confusion Matrix Visualization

```
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix # Import confusion matrix
```